

# Text as Data: Homework 1

Jeff Ziegler

August 16, 2017

*In this homework assignment we're going to analyze the first presidential debate from the 2012 election.*

## Problem 1

*To analyze the debate, we first need to load the debate and parse the content. On the coursewebsite, you'll find the file `debate1.html`. Download the file and open it in a browser. We will use `BeautifulSoup` to parse HTML file containing the debate transcript.*

- *Load the webpage into Python and use `BeautifulSoup` to create a searchable version of the debate. What tags can you use to identify statements?*

```
1 # import libraries
2 from bs4 import BeautifulSoup
3 from urllib import urlopen
4 import re
5 import os
6 import csv
7
8 # load .html file in gitHub folder (folder location will differ by user)
```

- *Note that not all of the statements contain information about the speaker. Devise a rule to assign the unlabeled statements to speakers. For substantive reasons, we would like to define a single statement as any uninterrupted speech from a candidate. We'll say a candidate is interrupted when the transcript says that a new speaker has begun. In other words, cross talk doesn't count as an interruption. Create a list with just the text (not the tags) of each statement as an element. Some statements are split among several tags; these will need to be concatenated according to the rule you devised above. Remember to filter out notes about audience behavior.*

```
1
2 # load .html file in gitHub folder (folder location will differ by user)
3 # find statements within <p> in HTML
4 pageText = BeautifulSoup(open('Documents/Git/WUSTL-textAnalysis/Debate1.
    html')).findAll('p')
5
```

```

6 # we know that there are three speakers
7 # Speakers: FORMER GOV. MITT ROMNEY, R-MASS; PRESIDENT BARACK OBAMA;
8 # JIM LEHRER, MODERATOR
9 # but even if we didn't know which names to search for, it appears
10 # they are labeled by all caps
11 # which is how we'll identify who is speaking and speaker changes
12
13 # create empty vector to be filled with statements
14 statements = []
15 # prior speaker is set to NULL, but will be filled with the most
16 # recent speaker
17 priorSpeaker = ''
18
19 # iterate over each text block (excluding the introduction and ending)
20 for i in pageText[6:477]:
21     # first, convert all <p> from bs4 object to strings to be searched
22     # and get rid of HTML in strings
23     # '\\' will appear, but it's just to escape the apostrophes
24     cleanedStatements = re.sub(re.compile('<.*?>'), '', str(i))
25     # then check if there is a fully capitalized word at the beginning
26     # of each statement
27     speakerLabelled = re.search('^[A-Z]+:', cleanedStatements)
28     # and if there is...
29     if speakerLabelled:
30         # record who the current speaker is (by checking which portion
31         # of the string matched the regex))
32         currentSpeaker = speakerLabelled.group()
33         # if the current speaker matches the prior speaker, add cleaned
34         # statement to the last full statement that was added
35         # and remove current speaker from every other statement
36         # except the first
37         if currentSpeaker == priorSpeaker:
38             # since no index is specified, .pop() removes and returns
39             # the last item in the list
40             # Note: there is an extra space added because otherwise
41             # append will crunch words together
42             statements.append(statements.pop() + " " +
43                             cleanedStatements.replace(currentSpeaker, ''))
44             # if the current speaker is different than prior speaker,
45             # add cleaned statement on its own
46         else:
47             statements.append(cleanedStatements)
48             # and reset prior speaker to the most recently recorded
49             speaker
50             priorSpeaker = speakerLabelled.group(0)
51             # if there is no speaker listed (does not match regex search),
52             # add cleaned statement to the last full statement that was added
53         else:
54             statements.append(statements.pop() + " " + cleanedStatements)

```

## Problem 2

Now we're going to do some more preprocessing to create a dataset that includes useful information about our texts. We will use a curated dictionary list from Neal Caren. The positive words are at <http://www.unc.edu/~ncaren/haphazard/positive.txt> and the negative words are at <http://www.unc.edu/~ncaren/haphazard/negative.txt>.

- Load the positive and negative words into python. Use the `porter`, `snowball` and `lancaster` stemmers from the `nlTK` package to create stemmed versions of the dictionaries.

```
1 ##### Problem 2
2
3 # create function to load sentimental dictionaries
4 def loadWords(type, stemmer):
5     # open url specifying positive or negative dictionary
6     url = urlopen('http://www.unc.edu/~ncaren/haphazard/' + type + '.txt').
        read()
7     # since they are in .txt files, we need to split each word
8     # create the unstemmed dictionary
9     unstemmedDict = url.split('\n')
10    # determine which stemmer should be used
11    # (1) Porter
12    if stemmer=="Porter":
13        # for each word in dictionary, stem
14        stemmedDict = [nlTK.stem.PorterStemmer().stem(word) for word in
            unstemmedDict]
15    # (2) Snowball
16    elif stemmer=="Snowball":
17        # for each word in dictionary, stem
18        stemmedDict = [nlTK.stem.SnowballStemmer('english').stem(word) for
            word in unstemmedDict]
19    # (3) Lancaster
20    elif stemmer=="Lancaster":
21        stemmedDict = [nlTK.stem.LancasterStemmer().stem(word) for word in
            unstemmedDict]
22    else:
23        stemmedDict = unstemmedDict
24    # return both stemmed and unstemmed dictionaries
25    # by using set() instead of keeping them as lists
26    # removes duplicates
27    return [unstemmedDict, set(stemmedDict)]
28
29 # get basic positive and negative, unstemmed dictionaries
30 positiveWords = loadWords('positive', stemmer="None").pop(0)
31 negativeWords = loadWords('negative', stemmer="None").pop(0)
32
33 # run dictionary acquisition and stemming function for all stemmers
34 # (1) Porter
35 stemmedPositivePorter = loadWords('positive', stemmer="Porter").pop(1)
36 stemmedNegativePorter = loadWords('negative', stemmer="Porter").pop(1)
```

```

37
38 # (2) Snowball
39 stemmedPositiveSnowball = loadWords('positive', stemmer="Snowball").pop
    (1)
40 stemmedNegativeSnowball = loadWords('negative', stemmer="Snowball").pop
    (1)

```

- Using the original and stemmed dictionaries, we're going to create a statement by statement data set of the speech. The data set should have the following columns:

- 1) Statement number (place in debate)
- 2) Speaker
- 3) Number of non-stop words spoken
- 4) Number of positive words
- 5) Number of negative words
- 6) Number of lancaster stemmed positive words
- 7) Number of lancaster stemmed negative words
- 8) Number of porter stemmed positive words
- 9) Number of porter stemmed negative words
- 10) Number of snowball stemmed positive words
- 11) Number of snowball stemmed negative words

To create the data set, create a set of nested dictionaries that map each statement in the list created in Problem 1 to the each of the attributes described above. To calculate the values for items 3 - 11 above, you'll need to do the following to each statement:

- Discard punctuation
- Remove capitalization
- Remove stop words with the list of words provided here:  
'<http://jmlr.org/papers/volume5/lewis04a/a11-smart-stop-list/english.stop>'
- Tokenize the words
- Apply each of the stemmers, determining which of the words appear in the corresponding stemmed dictionaries

Write your dataset as a .csv file and save it to a working directory. Turn it in with your homework.

```

1 # create function that will easily check how many words are in
2 # corresponding dictionary list
3 def wordCount(inputStatement, dictionaries):
4     return len([x for x in inputStatement if x in dictionaries])
5 # create function to pull necessary info from each statement
6 def statementInfo(statement, documentContent, count):
7     # first, need to discard punctuation
8     removedPunctuation = re.sub("\W", " ", i)
9     # capitalization
10    removedCaps = removedPunctuation.lower()
11    # and tokenization
12    reducedStatements = nltk.word_tokenize(removedCaps)
13
14    # append documentContent with relevant info
15    documentContent.append({
16    # add to statementIter
17    "statementNumber": count,
18    "speaker": re.search('^[A-Z]+', statement).group(),
19    # record the number of --- in statements w/ no punctuation, caps,
20    # and reduced tokens:
21    # non-stop words
22    "NstopWords": len([x for x in reducedStatements if x not in stop_words]),
23    # number of positive words
24    "NposWords": wordCount(reducedStatements, positiveWords),
25    # number of negative words
26    "NnegWords": wordCount(reducedStatements, negativeWords),
27    # number of words in each positive and negative using:
28    # (1) Porter stem
29    "NposPorter": wordCount([nltk.stem.PorterStemmer().stem(y) for y in
30    reducedStatements], stemmedPositivePorter),
31    "NnegPorter": wordCount([nltk.stem.PorterStemmer().stem(y) for y in
32    reducedStatements], stemmedNegativePorter),
33    # (2) Snowball stem
34    "NposSnowball": wordCount([nltk.stem.SnowballStemmer('english').stem(y)
35    for y in reducedStatements], stemmedPositiveSnowball),
36    "NnegSnowball": wordCount([nltk.stem.SnowballStemmer('english').stem(y)
37    for y in reducedStatements], stemmedNegativeSnowball),
38    # (3) Lancaster stem
39    "NposLancaster": wordCount([nltk.stem.LancasterStemmer().stem(y) for y in
40    reducedStatements], stemmedPositiveLancaster),
41    "NnegLancaster": wordCount([nltk.stem.LancasterStemmer().stem(y) for y in
42    reducedStatements], stemmedNegativeLancaster)})
43
44 # create empty list to fill with statement info
45 statementCharacteristics = []
46 # begin document iterations at 0
47 statementIter = 0
48 for i in statements:
49     # execute statementInfo function for each statement
50     # begin document iterations at 1
51     statementIter +=1

```

```

46 statementInfo(i, statementCharacteristics, count=statementIter)
47
48 # with data now assigned to dictionary
49 # write content to .csv
50 with open('Documents/Git/WUSTL_textAnalysis/statmentInfo.csv', 'wb') as f:
51     w = csv.DictWriter(f, fieldnames=("statementNumber", "speaker",
52     "NposWords", "NnegWords", "NposPorter", "NnegPorter",
53     "NposSnowball", "NnegSnowball", "NposLancaster", "NnegLancaster"))
54     w.writeheader()
55     for item in statementCharacteristics:
56         w.writerow(item)

```

### Problem 3

Using our new data set, let's make some observations about the debate

- Load the data into R
- Create a visualization that compares the overall positive and negative word rate for Obama, Romney, and Lehrer. What patterns do you notice? There is no one right answer, be creative!
- Using your data set, examine trends in each candidate's statements and Lehrer's speeches. Do you notice any
  - i) Trends in the measured tone?
  - ii) Response to the other candidate's tone (examining who spoke previously)?
  - iii) Overall interesting patterns? (this is an intentionally vague question)