

Text as Data: Homework 1

Jeff Ziegler

August 15, 2017

In this homework assignment we're going to analyze the first presidential debate from the 2012 election.

Problem 1

To analyze the debate, we first need to load the debate and parse the content. On the coursewebsite, you'll find the file `debate1.html`. Download the file and open it in a browser. We will use `BeautifulSoup` to parse HTML file containing the debate transcript.

- *Load the webpage into Python and use `BeautifulSoup` to create a searchable version of the debate. What tags can you use to identify statements?*

```
1 # import libraries
2 from bs4 import BeautifulSoup
3 from urllib import urlopen
4 import re
5 import os
6 import csv
7
8 # load .html file in gitHub folder (folder location will differ by user)
9 # find statements within <p> in HTML
10 pageText = BeautifulSoup(urlopen('Documents/Git/WUSTL-textAnalysis/Debate1.html')).findAll('p')
```

- *Note that not all of the statements contain information about the speaker. Devise a rule to assign the unlabeled statements to speakers. For substantive reasons, we would like to define a single statement as any uninterrupted speech from a candidate. We'll say a candidate is interrupted when the transcript says that a new speaker has begun. In other words, cross talk doesn't count as an interruption. Create a list with just the text (not the tags) of each statement as an element. Some statements are split among several tags; these will need to be concatenated according to the rule you devised above. Remember to filter out notes about audience behavior.*

```
1 # we know that there are three speakers
2 # Speakers: FORMER GOV. MITT ROMNEY, R-MASS; PRESIDENT BARACK OBAMA;
3 # JIM LEHRER, MODERATOR
```

```

4 # but even if we didn't know which names to search for, it appears
5 # they are labeled by all caps
6 # which is how we'll identify who is speaking and speaker changes
7
8 # create empty vector to be filled with statements
9 statements = []
10 # prior speaker is set to NULL, but will be filled with the most
11 # recent speaker
12 priorSpeaker = ''
13
14 # iterate over each text block (excluding the introduction and ending)
15 for i in pageText[6:477]:
16     # first, convert all <p> from bs4 object to strings to be searched
17     # and get rid of HTML in strings
18     # '\ ' will appear, but it's just to escape the apostrophes
19     cleanedStatements = re.sub(re.compile('<.*?>'), ' ', str(i))
20     # then check if there is a fully capitalized word at the beginning
21     # of each statement
22     speakerLabelled = re.search('^[A-Z]+:', cleanedStatements)
23     # and if there is...
24     if speakerLabelled:
25         # record who the current speaker is (by checking which portion
26         # of the string matched the regex))
27         currentSpeaker = speakerLabelled.group()
28         # if the current speaker matches the prior speaker, add cleaned
29         # statement to the last full statement that was added
30         # and remove current speaker from every other statement
31         # except the first
32         if currentSpeaker == priorSpeaker:
33             # since no index is specified, .pop() removes and returns
34             # the last item in the list
35             # Note: there is an extra space added because otherwise
36             # append will crunch words together
37             statements.append(statements.pop() + " " +
38                             cleanedStatements.replace(currentSpeaker, ''))
39         # if the current speaker is different than prior speaker,
40         # add cleaned statement on its own
41         else:
42             statements.append(cleanedStatements)
43             # and reset prior speaker to the most recently recorded
44             speaker
45             priorSpeaker = speakerLabelled.group(0)
46         # if there is no speaker listed (does not match regex search),
47         # add cleaned statement to the last full statement that was added
48         else:
49             statements.append(statements.pop() + " " + cleanedStatements)

```

Problem 2

Now we're going to do some more preprocessing to create a dataset that includes useful information about our texts. We will use a curated dictionary list from Neal Caren. The positive

words are at <http://www.unc.edu/~ncaren/haphazard/positive.txt> and the negative words are at <http://www.unc.edu/~ncaren/haphazard/negative.txt>.

- Load the positive and negative words into python. Use the porter, snowball and lancaster stemmers from the nltk package to create stemmed versions of the dictionaries.

```
1 # create function to load sentimental dictionaries
2 def loadWords(type, stemmer):
3     # open url specifying positive or negative dictionary
4     url = urlopen('http://www.unc.edu/~ncaren/haphazard/' + type + '.txt').
        read()
5     # since they are in .txt files , we need to split each word
6     # create the unstemmed dictionary
7     unstemmedDict = url.split('\n')
8     # determine which stemmer should be used
9     # (1) Porter
10    if stemmer=="Porter":
11        # for each word in dictionary , stem
12        stemmedDict = [nltk.stem.PorterStemmer().stem(word) for word in
            unstemmedDict]
13    # (2) Snowball
14    elif stemmer=="Snowball":
15        # for each word in dictionary , stem
16        stemmedDict = [nltk.stem.SnowballStemmer('english').stem(word) for
            word in unstemmedDict]
17    # (3) Lancaster
18    elif stemmer=="Lancaster":
19        stemmedDict = [nltk.stem.LancasterStemmer().stem(word) for word in
            unstemmedDict]
20    else:
21        stemmedDict = unstemmedDict
22    # return both stemmed and unstemmed dictionaries
23    return [unstemmedDict, stemmedDict]
24
25 # get basic positive and negative , unstemmed dictionaries
26 positiveWords = loadWords('positive', stemmer="None").pop(0)
27 negativeWords = loadWords('negative', stemmer="None").pop(0)
28
29 # run dictionary acquisition and stemming function for all stemmers
30 # (1) Porter
31 stemmedPositivePorter = loadWords('positive', stemmer="Porter").pop(1)
32 stemmedNegativePorter = loadWords('negative', stemmer="Porter").pop(1)
33
34 # (2) Snowball
35 stemmedPositiveSnowball = loadWords('positive', stemmer="Snowball").pop
    (1)
36 stemmedNegativeSnowball = loadWords('negative', stemmer="Snowball").pop
    (1)
37
38 # (3) Lancaster
```

```

39 stemmedPositiveLancaster = loadWords('positive', stemmer="Lancaster").pop
   (1)
40 stemmedNegativeLancaster = loadWords('negative', stemmer="Lancaster").pop
   (1)

```

- *Using the original and stemmed dictionaries, we're going to create a statement by statement data set of the speech. The data set should have the following columns:*

- 1) *Statement number (place in debate)*
- 2) *Speaker*
- 3) *Number of non-stop words spoken*
- 4) *Number of positive words*
- 5) *Number of negative words*
- 6) *Number of lancaster stemmed positive words*
- 7) *Number of lancaster stemmed negative words*
- 8) *Number of porter stemmed positive words*
- 9) *Number of porter stemmed negative words*
- 10) *Number of snowball stemmed positive words*
- 11) *Number of snowball stemmed negative words*

To create the data set, create a set of nested dictionaries that map each statement in the list created in Problem 1 to the each of the attributes described above. To calculate the values for items 3 - 11 above, you'll need to do the following to each statement:

- *Discard punctuation*
- *Remove capitalization*
- *Remove stop words with the list of words provided here:*
'<http://jmlr.org/papers/volume5/lewis04a/a11-smart-stop-list/english.stop>'
- *Tokenize the words*
- *Apply each of the stemmers, determining which of the words appear in the corresponding stemmed dictionaries*

Write your dataset as a .csv file and save it to a working directory. Turn it in with your homework.

```

1 # create function that will easily check how many words are in
2 # corresponding dictionary list
3 def wordCount(inputStatement, dictionaries):
4     return len([x for x in inputStatement if x in dictionaries])
5 # create function to pull necessary info from each statement

```

```

6 def statementInfo(statement, documentContent, count):
7     # first, need to discard punctuation
8     removedPunctuation = re.sub("\W", " ", i)
9     # capitalization
10    removedCaps = removedPunctuation.lower()
11    # and tokenization
12    reducedStatements = nltk.word_tokenize(removedCaps)
13
14    # append documentContent with relevant info
15    documentContent.append({
16    # add to statementIter
17    "statementNumber": count,
18    "speaker": re.search('^[A-Z]+', statement).group(),
19    # record the number of --- in statements w/ no punctuation, caps,
20    # and reduced tokens:
21    # non-stop words
22    "NstopWords": len([x for x in reducedStatements if x not in stop_words]),
23    # number of positive words
24    "NposWords": wordCount(reducedStatements, positiveWords),
25    # number of negative words
26    "NnegWords": wordCount(reducedStatements, negativeWords),
27    # number of words in each positive and negative using:
28    # (1) Porter stem
29    "NposPorter": wordCount([nltk.stem.PorterStemmer().stem(y) for y in
30    reducedStatements], stemmedPositivePorter),
31    "NnegPorter": wordCount([nltk.stem.PorterStemmer().stem(y) for y in
32    reducedStatements], stemmedNegativePorter),
33    # (2) Snowball stem
34    "NposSnowball": wordCount([nltk.stem.SnowballStemmer('english').stem(y)
35    for y in reducedStatements], stemmedPositiveSnowball),
36    "NnegSnowball": wordCount([nltk.stem.SnowballStemmer('english').stem(y)
37    for y in reducedStatements], stemmedNegativeSnowball),
38    # (3) Lancaster stem
39    "NposLancaster": wordCount([nltk.stem.LancasterStemmer().stem(y) for y in
40    reducedStatements], stemmedPositiveLancaster),
41    "NnegLancaster": wordCount([nltk.stem.LancasterStemmer().stem(y) for y in
42    reducedStatements], stemmedNegativeLancaster)})
43
44 # create empty list to fill with statement info
45 statementCharacteristics = []
46 # begin document iterations at 0
47 statementIter = 0
48 for i in statements:
49     # execute statementInfo function for each statement
50     # begin document iterations at 1
51     statementIter +=1
52     statementInfo(i, statementCharacteristics, count=statementIter)
53
54 # with data now assigned to dictionary
55 # write content to .csv
56 with open('Documents/Git/WUSTL_textAnalysis/statmentInfo.csv', 'wb') as f:

```

```

51 w = csv.DictWriter(f, fieldnames=("statementNumber", "speaker",
52 "NposWords", "NnegWords", "NposPorter", "NnegPorter",
53 "NposSnowball", "NnegSnowball", "NposLancaster", "NnegLancaster"))
54 w.writeheader()
55 for item in statementCharacteristics:
56     w.writerow(item)

```

Problem 3

Using our new data set, let's make some observations about the debate

- Load the data into R
- Create a visualization that compares the overall positive and negative word rate for Obama, Romney, and Lehrer. What patterns do you notice? There is no one right answer, be creative!
- Using your data set, examine trends in each candidate's statements and Lehrer's speeches. Do you notice any
 - i) Trends in the measured tone?
 - ii) Response to the other candidate's tone (examining who spoke previously)?
 - iii) Overall interesting patterns? (this is an intentionally vague question)